


# SINGLE-DIMENSIONAL OPTIMIZATION



# The minimization problem



# Specific objectives

- **Finding global minimum**

- Obtain the lowest possible value of the function
- A very hard problem to solve generally.

- **Finding local minimum**

- Smallest value with finite neighborhood
- Relative easier problem to solve

# A detour – local **root** finding problem

- Consider the problem of finding zeros for  $f(x)$
- Assume that you know
  - the point  $a$  where  $f(a)$  is positive
  - the point  $b$  where  $f(b)$  is positive
  - that  $f(x)$  is continuous between  $a$  and  $b$ .
- How would you proceed to find  $x$  such that  $f(x) = 0$ ?

# Bisection Method

```
## func : an arbitrary function
## lo    : x value where func(x) is negative
## hi    : x value where func(x) is positive
## e     : tolerance of errors in x
def binaryRoot(func, lo, hi, e):
    itr = 0          ## number of iteration
    while(True):    ## loop until returns
        d = hi-lo    ## range of x
        point = (hi+lo)/2.0  ## midpoint in x
        fpoint = func(point)  ## midpoint in y
        if ( fpoint < 0 ):    ## if y is negative..
            d = lo-point
            lo = point        ## set it as new lo
        else:
            d = point-hi
            hi = point        ## if not, set it as new hi
    ## return if range is small enough, or y is exactly zero
    if ( ( abs(d) < e ) or ( fpoint == 0 ) ):
        print("iteration:\t" + str(itr))
        print("point:\t" + str(point))
        print("fpoint:\t" + str(fpoint))
        print("d:\t" + str(d))
        return point;
    itr += 1
```

See at bios615\_2\_3.ipynb

# Evaluating the performance

```
import math  
binaryRoot(math.sin, 0-math.pi/math.sqrt(2), math.pi/2, 1e-5)
```

```
iteration:      18  
point:  1.44799433991e-06  
fpoint: 1.4479943399e-06  
d:      -7.23311957526e-06
```

# Improving the bisection method

- Approximation using linear interpolation

$$f^*(x) = f(a) + (x - a) \frac{f(b) - f(a)}{b - a}$$

- Root finding strategy
  - Select a new trial point such that  $f^*(x) = 0$

# False Position Method

```
def linearRoot(func, lo, hi, e):
    itr = 0          ## number of iteration
    flo = func(lo) ## need y values at the bounaries
    fhi = func(hi)
    while(True): ## loop until returns
        d = hi-lo          ## range of x
        point = lo + d * flo / (flo - fhi)
        fpoint = func(point) ## midpoint in y
        if ( fpoint < 0 ): ## if y is negative..
            d = lo-point
            lo = point          ## set it as new lo
            flo = fpoint        ## replace the bounary value, too
        else:
            d = point-hi
            hi = point ## if not, set it as new hi
            fhi = fpoint ## replace the bounary value, too
    ## return if range is small enough, or y is exactly zero
    if ( ( abs(d) < e ) or ( fpoint == 0 ) ):
        print("iteration:\t" + str(itr))
        print("point:\t" + str(point))
        print("fpoint:\t" + str(fpoint))
        print("d:\t" + str(d))
        return point;
    itr += 1
```

See at bios615\_2\_3.ipynb



# Comparing the performance

```
import math  
binaryRoot(math.sin, 0-math.pi/math.sqrt(2), math.pi/2, 1e-5)
```

```
iteration:      18  
point:  1.44799433991e-06  
fpoint: 1.4479943399e-06  
d:      -7.23311957526e-06
```

```
linearRoot(math.sin, 0-math.pi/math.sqrt(2), math.pi/2, 1e-5)
```

```
iteration:      6  
point:  2.58493941423e-26  
fpoint: 2.58493941423e-26  
d:      -2.63766959334e-20
```

# Comparing with R's **uniroot()**

```
> uniroot(sin, c(0-base::pi/sqrt(2),base::pi/2), tol=1e-5)
```

```
$root
```

```
[1] -1.376062e-11
```

```
$f.root
```

```
[1] -1.376062e-11
```

```
$iter
```

```
[1] 6
```

```
$init.it
```

```
[1] NA
```

```
$estim.prec
```

```
[1] 5e-06
```

# Summary - local root finding

- **Two simple methods**
  - Bisection method : `binaryRoot()`
  - False position method : `linearRoot()`
- **In the bisection method, the bracketing interval is halved at each step**
- **For well-behaved function, the false position method will converge faster, but there is no performance guarantee.**

# Back to the minimization problem

- Consider a complex function  $f(x)$  (e.g likelihood)
- Goal is to find  $x$  which  $f(x)$  is maximum or minimum value
- Maximization and minimization are equivalent
  - Replace  $f(x)$  with  $-f(x)$

# Notes from **root** finding

- **Two approaches possibly applicable to minimization problems:**
- **Bracketing**
  - Keep track of intervals containing solution
- **Accuracy**
  - Recognize that the solution has limited precision

# Considering the machine precision

- When estimating the minima and bracketing intervals, floating-point accuracy must be considered
- In general, if the machine precision is  $\epsilon$ , the achievable accuracy is no more than  $\sqrt{\epsilon}$
- In **double** precision floating-point number, it is ~8 decimal digits, since **double** has ~16 decimal digits of precision

# More **details** on machine precision

- $\sqrt{\epsilon}$  comes from the second-order Taylor approximation around a local minima

$$f(x) \approx f(b) + \frac{1}{2}f''(b)(x - b)^2$$

- For functions where higher order terms are important, accuracy could be even lower
  - For example, the minimum for  $f(x) = 1 + x^4$  is only estimated to about  $\epsilon^{\frac{1}{4}}$

# Steps for minimization strategy


1. Bracket minimum
  - Search for regions that contains local minima
2. Successively tighten bracket interval
  - .. using local minimization



# Detailed minimization strategy

- Find 3 points such that
  - $a < b < c$
  - $f(b) < f(a)$  and  $f(b) < f(c)$
- Then search for minimum by
  - Selecting trial points in the interval
  - Keep minimum values and flanking points

# Minimization **after** bracketing




# What is the best **location** of the new point X?




- Here, we cannot use a bisection method..

# What we **want** is..



- **We want to minimize the size of the next interval which will be either (A, X) or (B, C)**
  - If  $f(X) < f(B)$ , the next search interval will be (B, C)
  - If  $f(B) < f(X)$ , the next search interval will be (A, X)
- **What would be a principled strategy to select X?**

# Minimizing the **worst**-case damage




- The next interval will have length either  $1 - w$  or  $w + z$

- Optimal conditions are 
$$\begin{cases} 1 - w = w + z \\ \frac{z}{1 - w} = w \end{cases}$$

- Solving the equation leads to  $w = \frac{3 - \sqrt{5}}{2} = 0.38197$


# The golden search



# The initial bracketing **triplet**




# The golden ratio



The ratio between the new and old bracket size is  
1 : 1.618, which is known as golden ratio



# The golden ratio



The ratio between the new and old bracket size is  
1 : 1.618, which is known as golden ratio

# Summay - golden search

- Reduces bracketing by ~40%
- Performance is independent of the function that is being minimized
- In many cases, better schemes are available.

# Implementation of golden search

```
## goldenSearch() for general minimization
## func : an arbitrary function
## a : smallest value in bracket triplet
## b : middle value in bracket triplet
## c : largest value in bracket triplet
## e : relative precision tolerance
def goldenSearch(func, a, b, c, e):
    itr = 0
    fb = func(b)
    ## when b is very small, need a higher precision
    while( abs(c-a) > abs(b*e) ):
        if ( b > (a+c)/2 ): ## as illustrated in the lecture
            x = a + 0.38196 * (c-a)
        else:
            ## x is on the left side
            x = c - 0.38196 * (c-a)
        fx = func(x)      ## evaluate function
```

```

    if ( fx < fb ): ## if found a new minimum
        if (x > b):
            a = b
        else:
            c = b
        b = x
        fb = fx
    else:
        if ( x < b ):
            a = x
        else:
            c = x;
    itr += 1
print("itr:\t" + str(itr))
print("x:\t" + str(b))
print("y:\t" + str(fb))
print("d:\t" + str(c-a))
return b;

```

# A running example

```
def foo(x):  
    return 0-math.cos(x)
```

```
goldenSearch(foo, 0-math.pi/math.sqrt(2),math.pi/4,math.pi/2,1e-5)
```

```
itr:      69  
x:      -3.57189823006e-09  
y:      -1.0  
d:      2.32857231782e-14
```

# Using **existing** function in R

Global variable for tracking  
the # of function calls

```
> itr <- 0
> foo <- function(x) { itr <- itr +1 ; return (-cos(x)) }
> optimize(foo,c(-base::pi/sqrt(2),base::pi/2),tol=1e-5)
$minimum
[1] -3.995917e-07


$objective
[1] -1

> itr
[1] 9
```

# Can we **improve** the golden search?

- As with root finding, performance can improve substantially when interpolation methods are used
- However, a linear approximation won't work in this case, why?

# Approximation using parabola





# Parabolic interpolation

- Often converges faster than other algorithms
- However, there is no guarantee that this interpolation always works for an arbitrary function.
- Because golden search provides worst-case performance guarantee, it can be used as a fall-back for uncooperative functions.

# State-of-the-art : **Brent's** algorithm

- **Track 6 points (not all distinct)**
  - The bracket boundaries  $(a,b)$
  - The current minimum  $x$
  - The second and third smallest value  $(w, v)$
  - The new points to be examined  $u$
- **Use parabolic interpretation**
  - Using  $(x, w, v)$  to propose new value for  $u$ ,
  - Additional case is required to ensure  $u$  falls between  $a$  and  $b$
- **Implemented as in R functions**
  - **uniroot** for 1-dimensional root finding
  - **optimize** for 1-dimensional optimization

# Newton-Raphson method (1669)

- **Key idea**
  - Assume that the derivative function is available.
  - Use the derivative to find the next point with a linear interpolation
- **For root-finding problem, use the first derivative of  $f(x)$**

$$x_{n+1} = x_n - \frac{f(x_n)}{f'(x_n)}$$

- **For singled-dimensional optimization, use the second derivative**

$$x_{n+1} = x_n - \frac{f'(x_n)}{f''(x_n)}$$

# Properties of the Newton-Raphson method

- **Pros**

- Easy to implement
- Works for high-dimensional cases, too
- Quadratic convergence

- **Cons**

- Requires derivatives
- Convergence is not guaranteed
- Requires a big Jacobian or Hessian matrix for high-dimensional problems.
  - Quasi-newton method can be a fix

# Examples of the Newton-Raphson method

- Finding the reciprocal without division

$$f(x) = a - \frac{1}{x} \Rightarrow x_{n+1} = x_n - \frac{f(x_n)}{f'(x_n)} = x_n - \frac{a - \frac{1}{x_n}}{\frac{1}{x_n^2}} = x_n(2 - ax_n).$$

- Finding the square root with basic arithmetic operations

$$f(x) = x^2 - a \Rightarrow x_{n+1} = x_n - \frac{f(x_n)}{f'(x_n)} = x_n - \frac{x_n^2 - a}{2x_n} = \frac{x_n}{2} + \frac{a}{2x_n}.$$

# Summary

- **Root finding algorithms**
  - Bisection Method : Simple but likely less efficient
  - False Position Method : More efficient for well-behaved functions
  -
- **Single-dimensional minimization**
  - Golden Search: ~38% reduction of interval per iteration
  - Parabola Method: More efficient for well-behaved functions
  - Brent's Method: Combination of above two. State-of-the-art
  - Newton-Raphson Method: Quadratic convergence w/derivatives.